

CS336 Computer Networks - Project Specification Spring 2020

End-to-End Detection of Network Compression

Due: April 2nd, 2020 at 11:59PM

1 Project Outcomes

1. Implement two network applications (client/server and standalone) that detect the presence of network compression by end-hosts.
2. Verify and validate your compression detection applications.

2 Project Overview

For this project you will implement two network applications to detect whether network compression is present and locate the compression link on the network path. One application is a client/server application and the other application is a standalone application which works in an uncooperative environment. That is, the application detects network compression without requiring a special software being installed on the other host. This is inspired by the work, [End-to-End Detection of Compression of Traffic Flows by Intermediaries](#), which is recommended that you read in detail up to Section V.




Figure 1: A simple topology between two nodes with a compression link on the path.

3 Components

3.1 Compression Detection Client/Server Application

You will implement the network compression detection *only in the cooperative environment*. In summary, your network application is a client/server application where the sender sends two sets of n UDP packets back-to-back (called packet train), and the receiver records the arrival time between the first and last packet in the train (packets 0 and $n - 1$).¹ The first packet train consists of all packets of size ℓ bytes in payload, filled with all 0's, while the second packet train contains random sequence of bits. You can generate the random sequence of bits using `/dev/random`. If the difference in arrival time between the first and last packets of the two trains ($\Delta t_H - \Delta t_L$) is more than a fixed threshold $\tau = 100\text{ ms}$, the client application reports **Compression detected!**, whereas when the time difference is less than τ , there was probably no compression link on the path and the client application should display **No compression was detected**.

Every UDP packet in the packet train should have the following characteristics (Figure 2):

1. Don't Fragment flag in IP header set to 1
2. Source Port: 9876
3. Destination Port²: 8765

¹If tail packets are dropped, the application records the arrival time of the last packet received in the train instead.

²That requires the server to listen to this port to capture any incoming packets.

4. Packet ID: the first 16 bits of payload is reserved for a unique packet ID for each UDP packet in the train starting with packet ID 0.




Figure 2: UDP Packet Structure

Your client/server application consists of three phases:

1. **Pre-Probing Phase TCP Connection**

The client application initiates a TCP connection to the server. Upon successful connection, it then passes all the contents of the configuration file to the server (see Section 4). The server uses the provided information to detect whether compression is present or not. The TCP connection is released upon successful transmission of the contents of the configuration file.

2. **Probing Phase**

In this phase, the UDP packets are sent. The client, after sending n UDP packets with low entropy data, must wait Inter-Measurement Time (γ) seconds before sending the packet train with high entropy data. This is to ensure that the two probing phases will not interfere with one another.

3. **Post-Probing Phase TCP Connection**

After all packets are arrived, the server performs the necessary calculations required to detect the presence of compression on the path between the server and the client. After the probing phase, the client application initiates another TCP connection. Once the connection is established, the server sends its findings to the client.

3.2 (2) Compression Detection Standalone Application

This is a standalone application that detects network compression, without requiring any cooperation from the server, other than responding to standard network events. This application also locates the link where the compression is applied.

Using a raw socket implementation, your program sends a single *head* SYN packet. This is followed by a train of UDP packets, and a single *tail* SYN packet. The SYN packets are sent to two different ports (x and y) that are not among well-known ports. Sending SYN packets to closed ports should trigger RST packets from the server. The application only needs to record the arrival time of these two RST packets, and may ignore any ICMP messages generated by the UDP packets. The difference between the arrival time of the RST packets is then used to determine whether compression was detected or not. In event of either of the RST packets lost (you need a timer to infer loss) or the server not responding with an RST packet, your application will output **Failed to detect due to insufficient information.** and terminates.

The UDP packets are constructed in exact same way as described in Section 3.1, with one exception: their TTL field in the IP header are set manually to a fixed value for all UDP packets in the packet train. This information is provided by the user (see Section 4). The sequence of packets the standalone application generates is illustrated in Figure 3.




Figure 3: The illustration of the packet train generated by the standalone compression detection application.

4 Program Interface

You are to implement a Linux-based command-line program that takes one command-line argument: the *configuration file*.³ Below are examples of how one would execute your applications in Linux:

```
#On the client system
./compdetect_client myconfig.json

#On the server system
./compdetect_server myconfig.json

#On any system (the standalone program)
./compdetect myconfig.json
```

4.1 Configuration File

Your applications must use a configuration file to obtain the values for the parameters listed in this section. You may add other parameters as needed to this list. Depending on which of your three applications are reading from the configuration file, some parameters may be irrelevant, thus should be ignored. For instance, the server application will only use “Port Number for TCP (Pre-/Post- Probing Phases)” and ignore the rest. Or, only the standalone application uses the TTL value provided in the configuration file.

1. The Server’s IP Address
2. Source Port Number for UDP
3. Destination Port Number for UDP
4. Destination Port Number for TCP Head SYN, x
5. Destination Port Number for TCP Tail SYN, y
6. Port Number for TCP (Pre-/Post- Probing Phases)
7. The Size of the UDP Payload in the UDP Packet Train, ℓ (*default value: 1000B*)
8. Inter-Measurement Time, γ (*default value: 15 seconds*)
9. The Number of UDP Packets in the UDP Packet Train, n (*default value: 6000*)
10. TTL for the UDP Packets (*default value: 255*)

Good luck!

³You may choose any configuration file format as you wish. Some popular formats include but not limited to JSON, INI, and XML).